

Layout-Aware Translation Benchmark

A deterministic two-column document with a structured table

1. Introduction

Document translation quality depends on reliable reading order. The left column must remain separate from the right column, while headings and captions should provide context to the translation model.

2. Method

The standard pipeline parses native PDF text, detects semantic layout regions, recognizes table structure, and applies OCR only when required.

3. Results

Stable bounding boxes let LingoPane place translated text over the original page. A versioned document contract keeps the React interface independent from Docling internals.

Table 1. Expected pipeline responsibilities

Stage	Responsibility
PDF parser	Native text and coordinates
Layout	Semantic regions and order
Table	Rows, columns, and cells
oMLX	Local text translation